

# Computer Labs

## The Minix 3 Operating System And Virtual Box

2º L.EIC

Acknowledgements:

Pedro F. Souto (pfs@fe.up.pt)

# Goals

What is Minix 3?

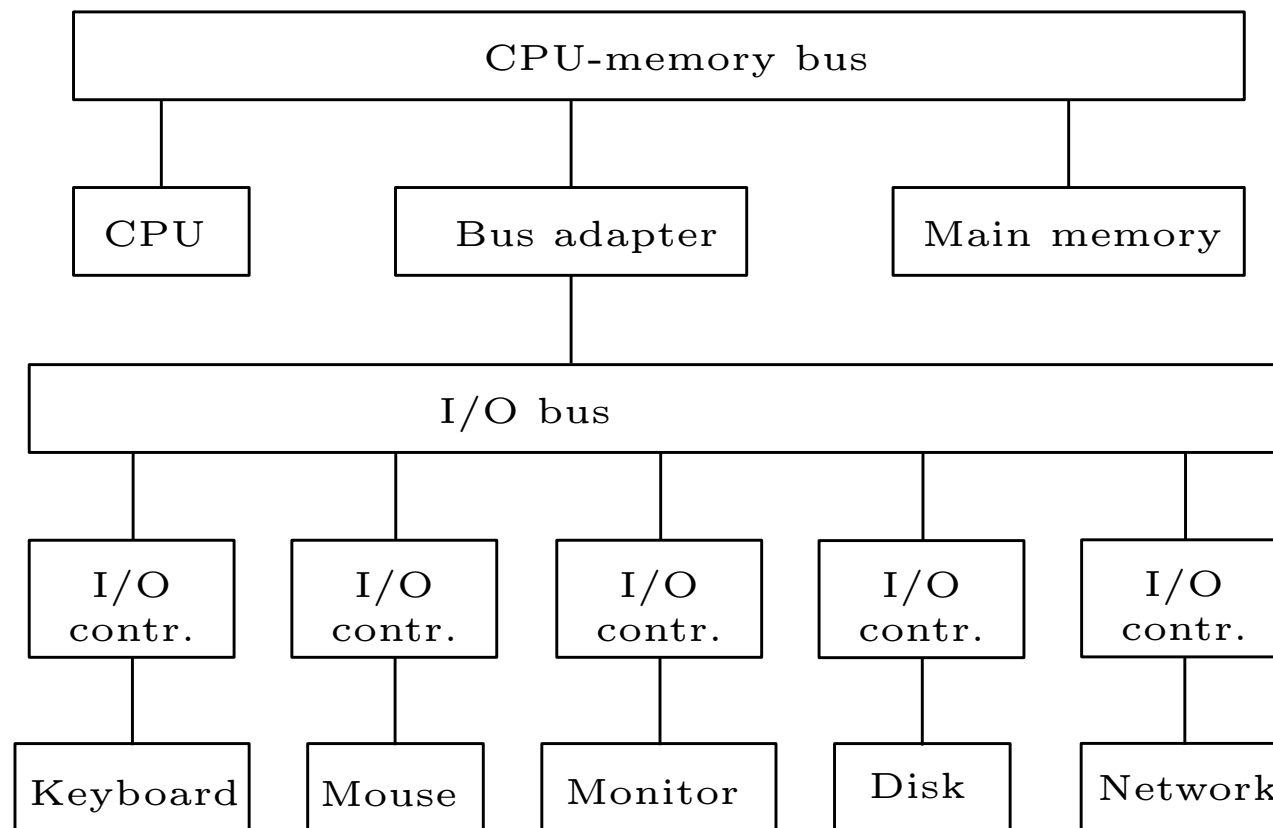
Why do we use Minix 3 in LCOM?

What is VirtualBox?

Why do we use VirtualBox in LCOM?

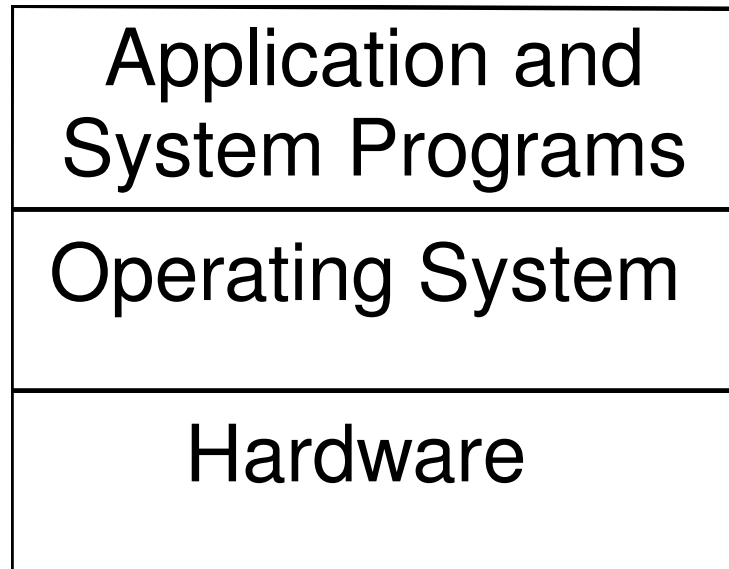
# LCOM Labs

- ▶ One of the goals of LCOM is that you learn to use the programmatic interface of the most common PC I/O devices



# Operating System (corrected)

- ▶ In most modern computer systems, access to the HW is mediated by the operating system (OS)



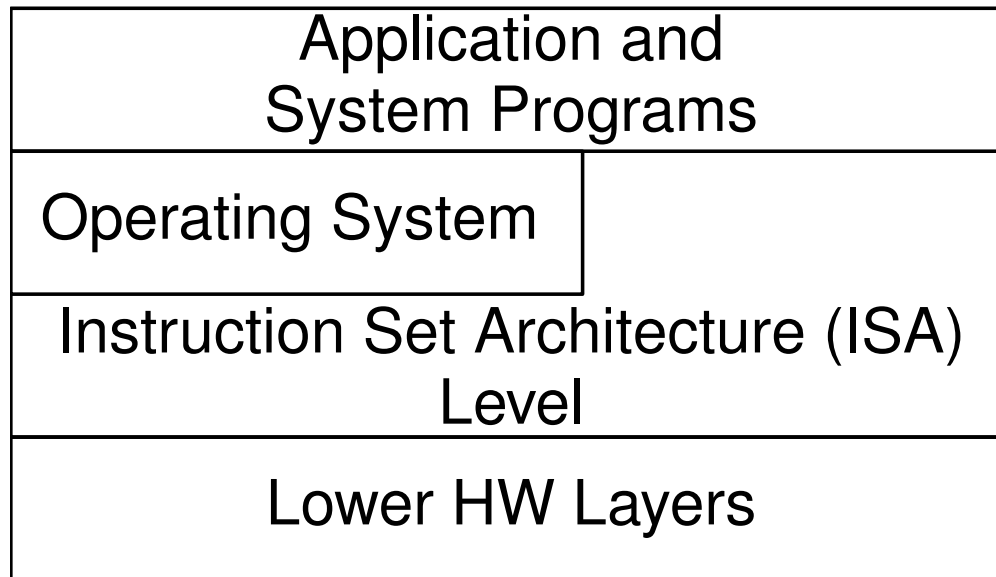
- ▶ I.e. user **processes** are not able to access directly the HW, mostly for reasons of:  
Reliability of the system  
Security

# Access to the HW-level Interface

Application and System Programs	
Operating System	
Instruction Set Architecture (ISA) Level	
Lower HW Layers	

- ▶ Most of the HW interface, actually the processor instruction set, is still available to user processes
- ▶ However, a few instructions are not directly accessible to user processes
  - ▶ Thus preventing user processes from interfering with:
    - Other processes most OSs are multi-process
    - The OS which manages the HW resources
- ▶ Instead, the operating system offers its own “instructions”, which are known as **system calls**.

# OS API: Its System Calls



**Extends** the ISA instructions with a set of “instructions”, system calls, that support concepts at a higher abstraction level

- ▶ OS system calls are too high-level for using directly the programmatic interface of I/O devices


**Hides** some ISA instructions

- ▶ The HW provides mechanisms that ensure that applications cannot bypass the OS API

**Issue** The OS API (of main-stream OSs) do not allow us to directly access the programmatic interface of I/O devices

# Parenthesis: OS vs. Kernel

- ▶ Usually, when we mention the OS we really mean the kernel
- ▶ An OS has several components



**Kernel** Which implements the system calls and manages the hardware

**Library** Which provides an API so that processes can make system calls

**Utilities** A set of “basic” programs, that allows a “user” to use the computing system resources

# Parenthesis: Layered Structure

- ▶ Structure typically used to address complex problems
  - ▶ It allows us to think about the **what** without worrying about the **how** (this is usually called **abstraction**)
- ▶ This has several advantages
  - Decomposition** An “intractable” problem is decomposed in smaller problems that can be solved
  - Modularity** Facilitates adding new functionality or changing the implementation, as long as the **interfaces are preserved**
- ▶ Your project will be a somewhat complex piece of code
  - ▶ To structure it in several layers may be very important for your success

Other SW layers			
Video Driver	Keyboard Driver	Timer Driver	Mouse Driver



# How is an OS/Kernel implemented?

**Monolithic** All OS services, e.g. file system or device drivers, are implemented at kernel level by the kernel

- ▶ Usually, the kernel is developed in a modular fashion
- ▶ However, there are no mechanisms that prevent one module from accessing the code, or even the data, of another module

**Micro-kernel** Most OS services are implemented as modules that execute in their own address spaces

- ▶ A module cannot access directly data or even code of another module
- ▶ There is however the need for some functionality to be implemented at kernel level, but this is minimal (hence the name)

# Monolithic Implementation

- ▶ Virtually all “main stream” OSs use this architecture
- ▶ It has lower overheads, and is faster

But is less reliable, because components are not isolated from each other

- ▶ If we used Linux or Windows instead of Minix, a bug in your program could just crash the entire system
  - ▶ This would make the development process very painful

# Minix 3: Micro-kernel Based

- ▶ It has a very small size kernel (about 6 K lines of code, most of it C)
- ▶ Most of the OS functionality is provided by a set of **privileged user-level** processes:

**Services** E.g. file system, process manager, VM server, Internet server, and the resurrection server.

**Device Drivers** All of them are user-level processes

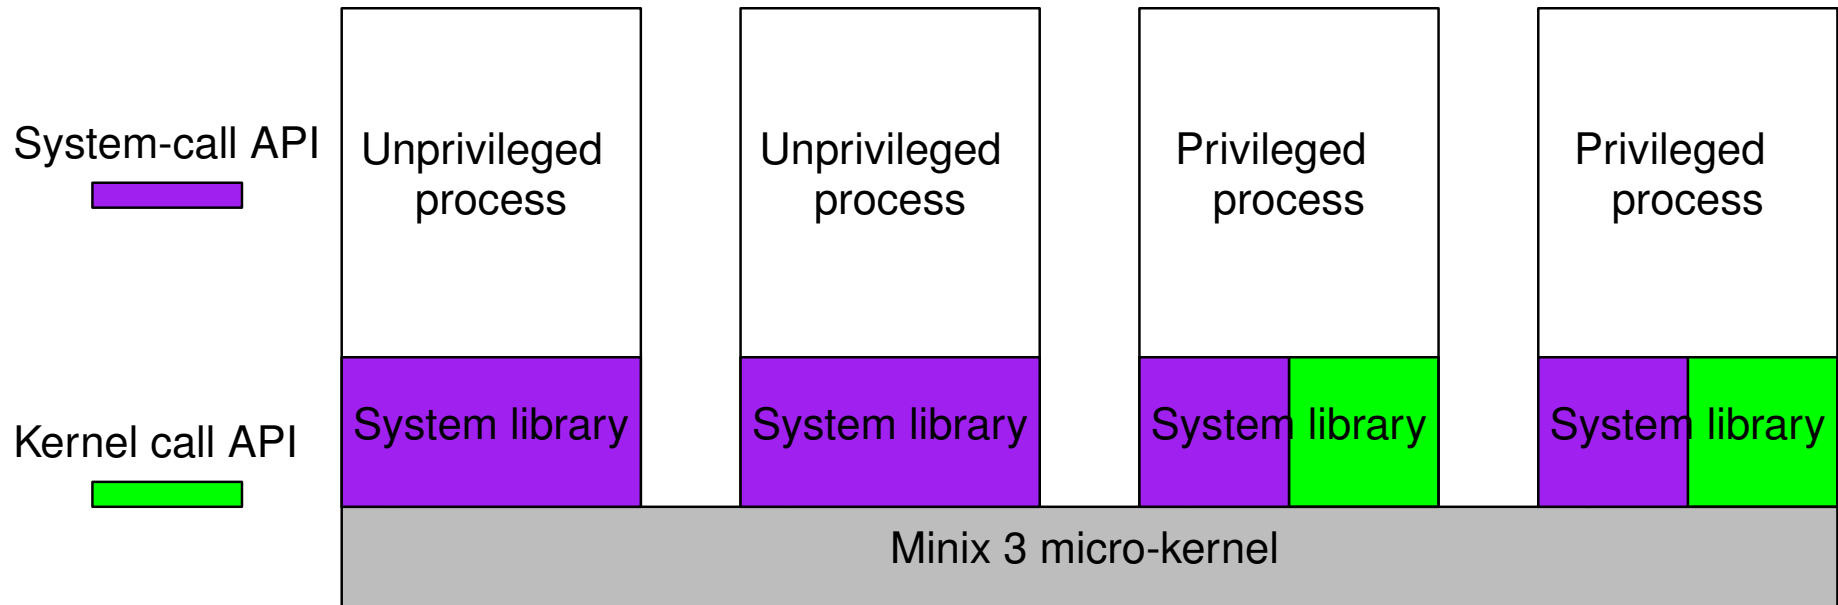
- ▶ Minix 3 provides an API, which is known as **kernel-calls**, that allow privileged processes to execute instructions required by device-drivers

- ▶ E.g. `sys_inb()`, which you'll use in Lab 2

**Note** Kernel-calls are (conceptually) different from system calls

- ▶ Any process can execute a system call
  - ▶ Only privileged processes are allowed to execute a kernel call

# Minix 3: Non-Privileged vs. Privileged User Processes



**Conclusion** by using Minix 3, LCOM processes not running in kernel mode can use the programmatic interface of I/O devices

- ▶ The development process is much less painful
  - ▶ Your processes do not belong to the kernel
  - ▶ Their actions can be controlled

Thus, bugs are much less harmful

# VirtualBox (1/3)

**Problem** Direct access to the programmatic interface of a computer's I/O devices raises **security risks**. E.g.

- ▶ If you are able to directly access a HDD (or an SSD), you can have access to the data it stores
  - ▶ Worse, you can install malware that can, e.g., spy on users, even long after the last time you used that computer
- ▶ In FEUP, these risks are unacceptable

**Solution** Use a **virtual machine**, VirtualBox in the case of LCOM

# VirtualBox (2/3)

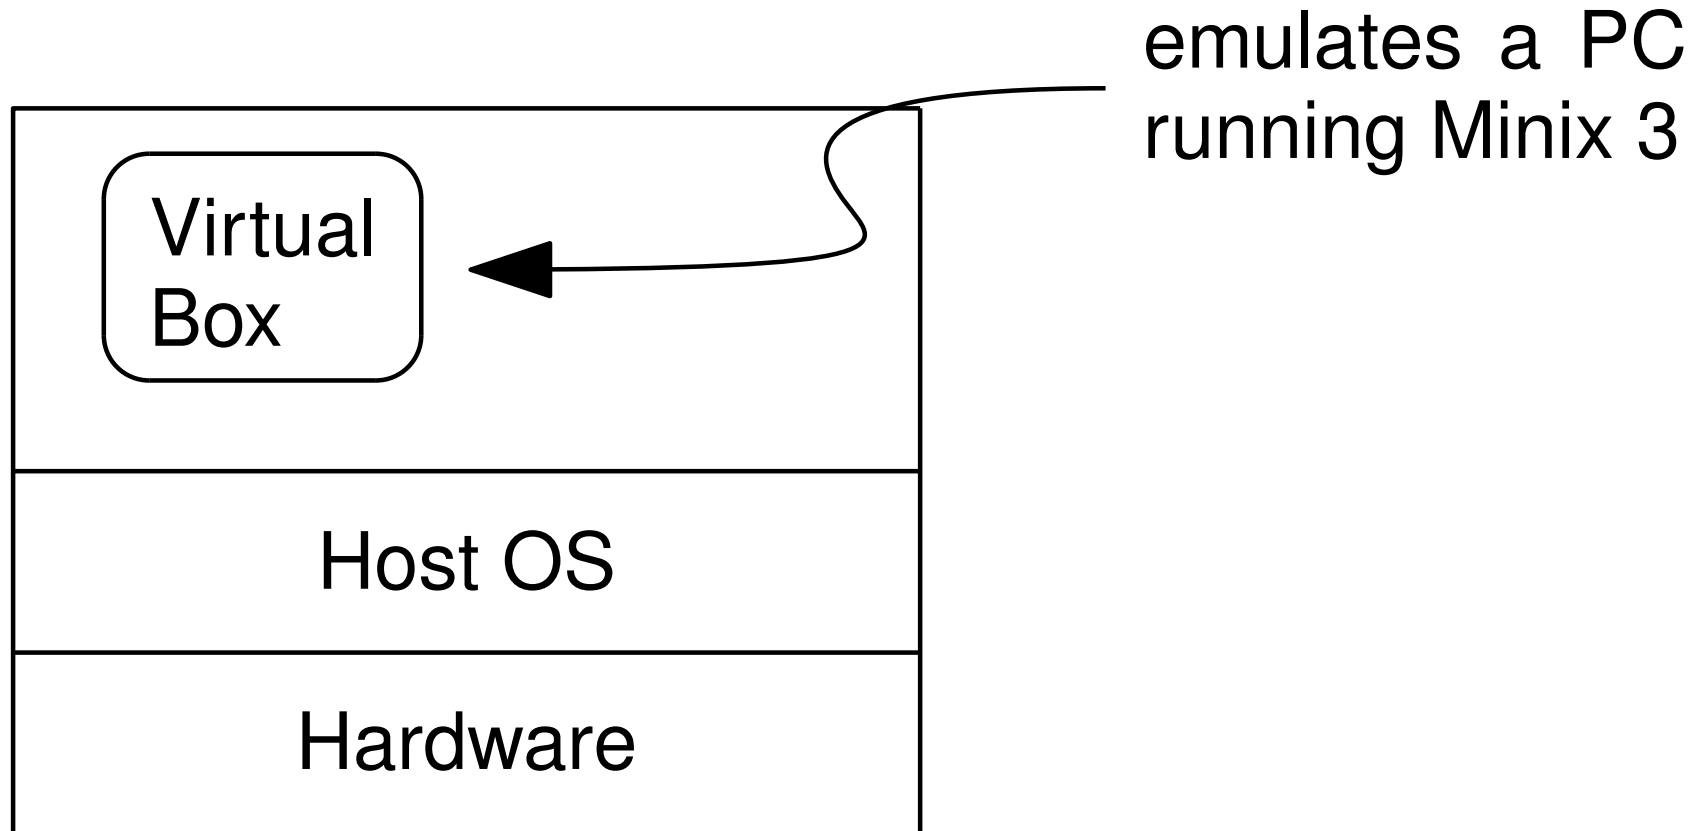
**What is VirtualBox?** Is a (system) virtual machine, more specifically:

- ▶ It is a program that emulates the HW of a PC
  - ▶ VirtualBox runs as a privileged process in a computer system
- ▶ It can run (most) programs that run on a PC without any modification
  - ▶ Both the operating systems
  - ▶ And applications or user programs.

**Why is this useful?** When you run a program in Virtual Box you can only access emulated resources not the physical resources. E.g.:

- ▶ Access to an emulated HDD (or SSD) exposes only the data stored in that emulated HDD (usually a file in a physical HDD), not the data in the entire (physical) HDD of **host** computer, i.e. the computer that runs the VM

## VirtualBox (3/3)



# Computer Labs: I/O Devices

## 2º L.EIC

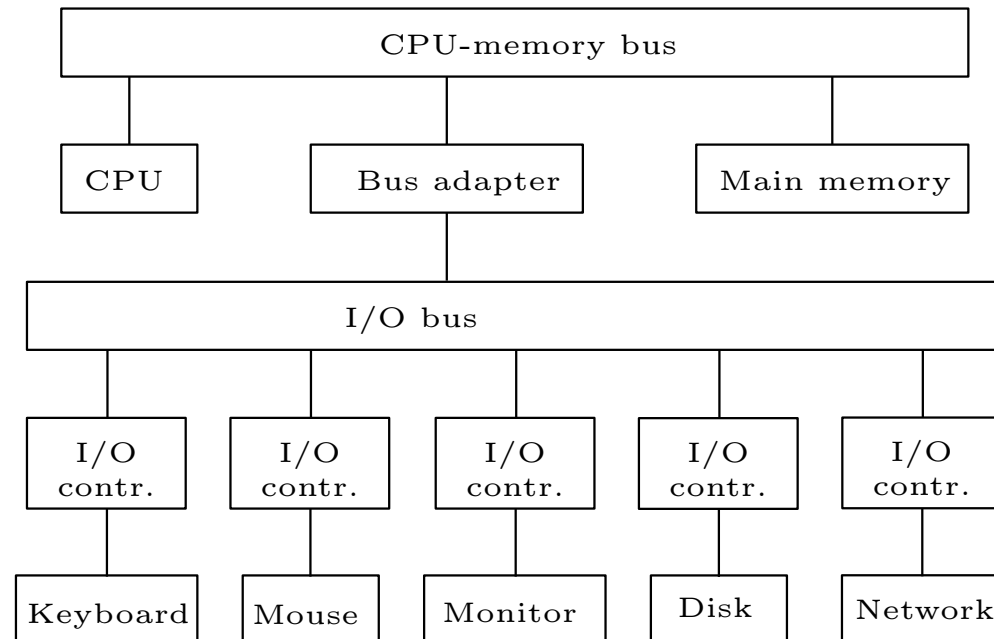
Acknowledgements:

Pedro F. Souto (pfs@fe.up.pt)



# I/O Devices

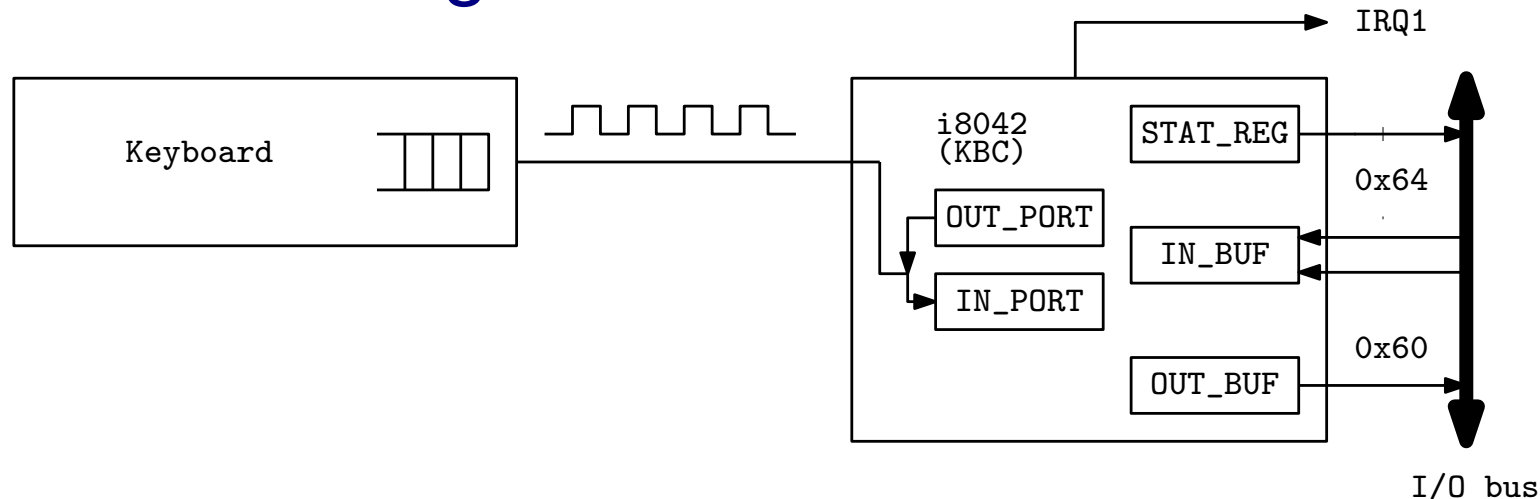
- ▶ I/O devices provide the interface between the CPU and the outside world.



# I/O Controllers

- ▶ Each I/O device is controlled by an electronic component, usually called **controller** or **adapter**.
- ▶ I/O controllers typically include three kinds of registers:
  - Control:** used to request I/O operations
  - Status:** used to get the state of the device or pending I/O operations
  - Data:** used to transfer data to/from the I/O devices
    - Input Data
    - Output Data
- ▶ Programming at the register level may require a detailed knowledge of the device's operation

# The KBC Registers



- ▶ The KBC has two registers at port `0x60`;  
**Input Buffer** (`IN_BUF`) used for sending commands to the keyboard (KBD commands)
  - ▶ The names are from the point of view of the KBC, not the CPU**Output Buffer** used for receiving scancodes and ...
- ▶ And two registers at port `0x64`  
**Status Register** for reading the KBC state  
**Not named** for writing KBC commands
  - ▶ Apparently, this is not different from the `IN_BUF` at port `0x60`
  - ▶ The value of input line `A2` is used by the KBC to distinguish KBC commands from KBD commands the `IN_BUF`

# How does the CPU access an I/O controller?

## Via memory-mapped I/O

- ▶ Portions of the address-space are assigned to I/O devices
- ▶ Access to an I/O device is done using the CPU's memory access instructions
- ▶ Can be used with any processor architecture

## Special I/O instructions

- ▶ I/O uses different address-space and each I/O device is assigned a portion of that address space
- ▶ CPU must provide special instructions to access the I/O address-space (I/O instructions)
  - ▶ The Intel CPU's used in the PC have always provided them
  - ▶ The ARM processors do not
- ▶ I/O instructions are legal only when executing at a high privilege level, typically that of the kernel/supervisor mode

# Intel's I/O Instructions

**Port** Is an abstraction of a device's controller register

- ▶ In the Intel documentation, a port is the name of an address in the I/O address space
- ▶ The I/O address space uses 16-bit addresses
- ▶ Two/four-consecutive 8-bit ports can be treated as 16/32-bit ports – should align them for performance

**Instruction IN** Input from port, i.e. read from an I/O register

- ▶ The source operand, i.e. the I/O port, is either a “byte immediate” or the DX register (a 16-bit register)
- ▶ The destination operand is one of the AL, AX and EAX registers, depending on the size of the port being accessed

```
in      al,    80h           ; read byte from port 80h
```

**Instruction OUT** Output to Port, i.e. write to an I/O register

```
mov     dx,    3F8h
out     dx,    al           ; write byte to port 3F8h
```

# Access to I/O Registers in C

**Issue** C does not provide any instruction for executing IN/OUT assembly instructions

**Solution** Use Minix 3 `SYS_DEVIO` kernel call for doing I/O

```
#include <minix/syslib.h>
```

```
int sys_inb(int port, u32_t *byte);
```

```
int sys_outb(int port, u32_t byte);
```

**Note** that the second argument of `sys_inb()` must be the address of a 32-bit unsigned integer variable.

**Hint** implement

```
util_sys_inb(int port, u8_t *byte)
```

- ▶ This is a wrapper to `sys_inb()`
- ▶ You can use it thereafter instead of `sys_inb()`

# PC's I/O address map

