

# Computer Labs: I/O and Interrupts

## 2º L.EIC

Pedro F. Souto (`pfs@fe.up.pt`)

February 27, 2025

# I/O Operation

- ▶ I/O devices are the interface between the computer and its environment
- ▶ Most of the time, the processor is not synchronized with its environment
  - ▶ I/O operations are **asynchronous** wrt the processor operation
- ▶ Usually, I/O devices are much slower than the processor
  - ▶ The processor **must wait** for an I/O device to complete its current operation before it can request a new one

# How Does the Processor Know about an I/O event?

**Polling** The processor polls the I/O device, i.e. reads a status register, to find out

**Interrupts** The I/O device notifies the processor, via the interrupt mechanism

# How Does the Processor Know about an I/O event?

**Polling** The processor polls the I/O device, i.e. reads a status register, to find out

**Response time** Highly variable – depends on what the processor has to do between consecutive polls.

**Interrupts** The I/O device notifies the processor, via the interrupt mechanism

**Response time** Usually responsive – depends on the time:

- ▶ interrupts are disabled or
- ▶ higher priority interrupts take to be served

# How Does the Processor Know about an I/O event?

**Polling** The processor polls the I/O device, i.e. reads a status register, to find out

**Response time** Highly variable – depends on what the processor has to do between consecutive polls.

**Efficiency/Overhead** Depends on the frequency of the event

- ▶ The more frequent the more efficient
- ▶ Assuming, polling at a constant rate

**Interrupts** The I/O device notifies the processor, via the interrupt mechanism

**Response time** Usually responsive – depends on the time:

- ▶ interrupts are disabled or
- ▶ higher priority interrupts take to be served

**Efficiency/Overhead** Depends on the frequency of the event

- ▶ The more frequent the less efficient
- ▶ Overhead per interrupt is higher than that per poll

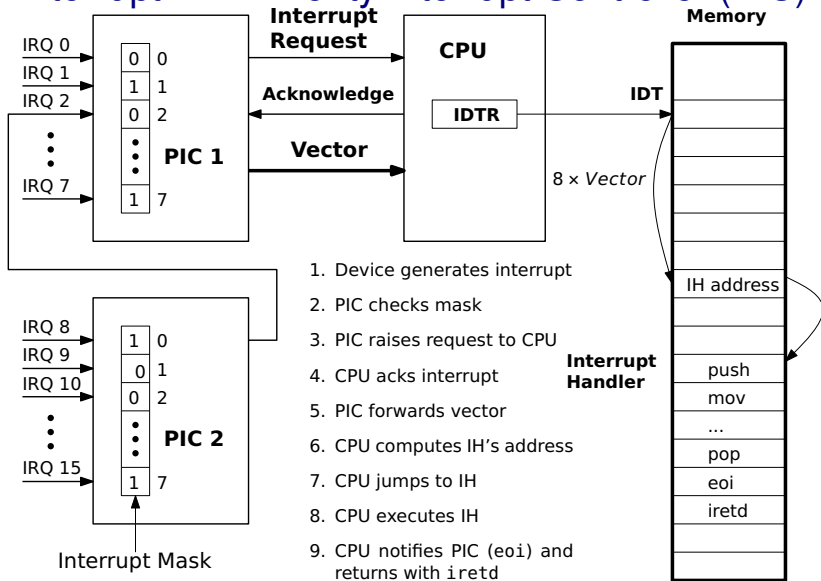
## Lab 2: `timer_test_int()`

**What to do?** Print one message per second, for a time interval whose duration is specified in its argument, by using:

- ▶ Timer 0 interrupts
- ▶ LCF function:

```
void timer_print_elapsed_time()
```

# PC Interrupt HW: Priority Interrupt Controller (PIC)



**Imp:** If a bit of the Interrupt Mask is set, the corresponding IRQ is disabled.

# PC Interrupts: IRQ Lines and Vectors

PIC 1	PIC 2	Device	Vector
IRQ0		Timer	0x08
IRQ1		Keyboard	0x09
IRQ2		PIC2	0x0A
	IRQ0	Real Time Clock	0x70
	IRQ1	Replace IRQ2	0x71
	IRQ2- IRQ7	Reserved	0x72-0x77
IRQ3		Serial port COM2	0x0B
IRQ4		Serial port COM1	0x0C
IRQ5		Reserved/Sound card	0x0D
IRQ6		Floppy disk	0x0E
IRQ7		Parallel port	0x0F

**IRQ line** Determined by the HW designer (IBM)

**Vector** Specified also by IBM, but can be configured at boot time. All that is needed is:

1. Configure the PIC
2. Configure the IDT (Interrupt Descriptor Table)



# Interrupt Handlers (IH)

- ▶ IHs are executed by the HW upon an interrupt
  - ▶ They run **asynchronously** wrt other code
  - ▶ They take no arguments
  - ▶ They return no values
- ▶ IHs used to be written in assembly
  - ▶ Need to perform I/O operations

```
isr_name:
    push ..                ; save all registers used
    ...                    ; IH instructions
    mov al, EOI             ; signal EOI
    out PIC1_CMD, al        ; to PIC1
    pop ...                 ; restore all registers used
    iretd
```

- ▶ But nowadays, they are usually written in C (for reasons of portability)

**Terminology** Interrupt handlers are also called interrupt service routines (ISR) and are part of the respective **device driver**

# Interrupt Handling in Minix 3

- ▶ In Minix, device drivers are implemented as **user-level processes**, rather than at the kernel-level
  - ▶ This was an important design decision in Minix 3

# Interrupt Handling in Minix 3

- ▶ In Minix, device drivers are implemented as **user-level processes**, rather than at the kernel-level
  - ▶ This was an important design decision in Minix 3

**Issue** How do you do interrupt handling?

- ▶ Interrupt handling requires performing operations that usually require special privileges

# Interrupt Handling in Minix 3

- ▶ In Minix, device drivers are implemented as **user-level processes**, rather than at the kernel-level
  - ▶ This was an important design decision in Minix 3

**Issue** How do you do interrupt handling?

- ▶ Interrupt handling requires performing operations that usually require special privileges

**Solution**

1. Perform only the bare minimum in the kernel: this is done by the **generic interrupt handler** (GIH)
2. Device specific operations are performed by the device drivers themselves at user level
  - ▶ Using kernel calls to perform privileged operations

# Minix 3: The Generic Interrupt Handler (GIH)

- Upon an interrupt, the GIH:
  1. Masks, in the PIC, the respective IRQ line.
  2. Notifies all the device drivers (DD) **interested** in that interrupt
  3. If possible, unmask, in the PIC, the respective IRQ line.
  4. Acknowledges the interrupt by issuing the `EOI` command to the PIC.
  5. Issues the `IRETD` instruction

**Issue 1** How does the GIH know that a DD is interested in an interrupt?

**Issue 2** How does the GIH notify a DD?

**Issue 3** How does a DD receive the notification of the GIH?

**Issue 4** How does the GIH know if it can unmask the IRQ line in the PIC?

**Issue 5** If the GIH does not unmask the IRQ line in the PIC, when, how and whom does it?

# Issue 1

How does the GIH know that a DD is interested in an interrupt?

# Issue 1

How does the GIH know that a DD is interested in an interrupt?

**Answer** The DD tells it, using kernel call:

```
int sys_irqsetpolicy(int irq_line, int policy, int *hook_id)
```

where

`irq_line` is the IRQ line of the device

`policy` use `IRQ_REENABLE` to inform the GIH that it can unmask the IRQ line in the PIC.

- ▶ This answers Issue 4: How does the GIH know if it can unmask the IRQ line in the PIC?

`hook_id` is both:

**input** an id to be used by the kernel on interrupt notification

**output** an id to be used by the DD in other kernel calls on this interrupt

- ▶ `sys_irqsetpolicy()` can be viewed as an interrupt notification subscription

## Minix 3: Other Interrupt Related Kernel Calls

`sys_irqrmpolicy(int *hook_id)` Cancels a previous interrupt notification subscription, by specifying a pointer to the `hook_id` returned by the kernel in `sys_irqsetpolicy()`

`sys_irgenable(int *hook_id)` Unmasks at the PIC an interrupt line associated with a previously subscribed interrupt notification, by specifying a pointer to the `hook_id` returned by the kernel in `sys_irqsetpolicy()`

`sys_irqdisable(int *hook_id)` Masks at the PIC an interrupt line associated with a previously subscribed interrupt notification, by specifying a pointer to the `hook_id` returned by the kernel in `sys_irqsetpolicy()`



## Issue 2

How does the GIH notify the DD of the occurrence of an interrupt?

## Issue 2

How does the GIH notify the DD of the occurrence of an interrupt?

**Answer** It uses the standard interprocess communication (IPC) mechanism used for communication:

- ▶ between processes;
- ▶ between the (micro) kernel and a process

More specifically, it uses **notifications**

**Minix 3 IPC** This is essentially a message based mechanism

- ▶ Processes send and receive messages to communicate with one another, and with the kernel.
- ▶ A **notification** is a special kind of message, used by the kernel to unsolicited communication with a user-level process.

## Issue 3 (1/2)

How does the DD receive the notification of the GIH?

## Issue 3 (1/2)

How does the DD receive the notification of the GIH?

**Short Answer** Just use the IPC mechanism.

**Useful Answer** Use some library calls provided by the  
`libdrivers` library

```

1: #include <lcom/lcf.h>
2: int ipc_status;
3: message msg;
4: while( 1 ) { /* You may want to use a different condition */
5:     /* Get a request message. */
6:     if( (r = driver_receive(ANY, &msg, &ipc_status)) != 0 ) {
7:         printf("driver_receive failed with: %d", r);
8:         continue;
9:     }
10:    if (is_ipc_notify(ipc_status)) { /* received notification */
11:        switch (_ENDPOINT_P(msg.m_source)) {
12:            case HARDWARE: /* hardware interrupt notification */
13:                if (msg.m_notify.interrupts & irq_set) { /* subscribed */
14:                    ... /* process it */
15:                }
16:                break;
17:            default:
18:                break; /* no other notifications expected: do nothing */
19:        }
20:    } else { /* received a standard message, not a notification */
21:        /* no standard messages expected: do nothing */
22:    }
23: }

```

Why: `msg.m_notify.interrupts`?

- ▶ Interrupt handlers take no arguments (and return no values)

Answer True, but usually an IH knows which interrupt request it is handling

- ▶ Minix 3 allows a DD to subscribe notifications on several interrupt lines

What is its value?

Answer It is based on the input value of `hook_id` passed by the DD in the corresponding `sys_irqsetpolicy()`.

- ▶ If a given interrupt is pending then the corresponding `hook_id` bit of `msg.m_notify.interrupts` is set.
- ▶ Why not just the `hook_id`?

What should `irq_set` value be?

- ▶ `irq_set` is used as a mask to test which interrupts are pending

## Issue 3 (2/2)

**Key Observation** In Minix 3, a DD is an event driven service that receives and processes messages

- ▶ either interrupt notifications from the kernel (GIH)
- ▶ or service requests from other processes

However, the programs in LCOM are not DD: they do not receive requests from other processes

## Lab 2: `timer_test_int()`

**What to do?** Print one message per second, for a time interval whose duration is specified in its argument.

1. Subscribe Timer 0 interrupts
2. Print message at 1 second intervals, by calling the LCF function:

```
void timer_print_elapsed_time()
```

3. Unsubscribe Timer 0 at the end

**How to design it?** It is not easy to come up with an API that can be used in the project

- ▶ Implement `int timer_subscribe_int()` to hide from other code i8254 related details, such as the IRQ line used
  - ▶ It returns, via its arguments, the bit number, that will be set in `msg.m_notify.interrupts` upon a TIMER 0 interrupt
- ▶ Implement the interrupt handler also in `timer.c`
- ▶ Implement the “interrupt loop” in `timer_test_int()`



## Issue 5 (and Last)

What if the GIH does not unmask the IRQ line in the PIC?

## Issue 5 (and Last)

What if the GH does not unmask the IRQ line in the PIC?

- ▶ I.e., if a DD does not set the `IRQ_REENABLE` policy in its interrupt subscription request (`sys_irqsetpolicy()`)

**Answer** The DD will have to do it, as soon as possible

- ▶ In most cases, you'll want to set the `IRQ_REENABLE` policy
  - ▶ In Lab 2, certainly

How can a DD unmask the IRQ line in the PIC??

- ▶ By calling `sys_irqenable(int *hook_id)`
  - ▶ Note that here `hook_id` should point to a variable with the value returned by the kernel in `sys_irqsetpolicy()`

That is, the kernel will unmask the IRQ line, upon request of the DD.

# Minix 3: Interrupt Sharing

- ▶ Minix 3 already includes its own Timer 0 IH
- ▶ By subscribing interrupts on IRQ line 0, the IH of your driver will not replace the IH of the kernel
  - ▶ Upon an interrupt generated by Timer 0, the kernel:
    1. executes its own IH, and
    2. notifies your driver
- ▶ This behavior stems from the need to share the interrupt lines among devices
  - ▶ In systems with the PIC (i8259), there are only 15 interrupt lines available
  - ▶ And many of them are actually hardwired, e.g. IRQ 0, which means that they cannot be shared among devices

**IMP** Using two IH for the same device is seldom what you want

- ▶ But is just what we need for Lab 2.

# Further Reading

- ▶ Lab 2 Handout, Section 4, The PC's Interrupt Hardware
- ▶ 8259A- Interrupt Priority Controller- Data Sheet, by Intel
- ▶ Using Interrupts
- ▶ Lab 2 Handout, Subsection 5.2 (Minix 3) Interrupt Handling